

# Running the Intelligent Document Insight Engine

The project is shipped as a self-contained Docker solution; you **don't need to install any Python libraries or set up a local runtime**. Follow the steps below to build, run, and interact with both Round 1A (outline extraction) and Round 1B (persona-aware ranking).

## 1. Prerequisites

Component	Required version	Notes
Docker Engine / Desktop	20.10 +	Linux, macOS, or Windows 10/11 with WSL 2
CPU & RAM	≥ 4 vCPU & 8 GB	Model weights load into memory during Round 1B
Disk space	≈ 1 GB	Image size after build is ≈ 700 MB

**Windows users:** enable WSL 2 and file-sharing for the drive where your `input/` and `output/` folders reside [12](#).

## 2. Obtain the code package

Download `connecting-dots.zip` from the previous assistant message.

Unzip to a working directory, e.g. `C:\doc-engine` or `~/doc-engine`.

Inside you will see:

```
app/  
Dockerfile  
docker-compose.yml  
README.md  
input/  
output/
```

## 3. Build the Docker image

Open a terminal in the project root and run:

```
docker build -t connecting-dots .
```

*First build takes 3 min on a typical laptop; subsequent builds are cached.*

## 4. Prepare test PDFs

Copy one or more PDF files (≤ 50 pages each) into the `input/` folder.

Ensure filenames contain no spaces; the container scans `*.pdf`.

## 5. Run Round 1A – outline extraction

```
docker run -- rm ^
  -v $(pwd)/input:/app/input ^
  -v $(pwd)/output:/app/output ^
  connecting-dots --mode 1A
```

On Windows PowerShell replace `$(pwd)` with `${PWD}`.

Results: for every `document.pdf` you will find `document_outline.json` in `output/` containing the hierarchical heading list.

## 6. Run Round 1B – persona-aware ranking

```
docker run -- rm ^
  -v $(pwd)/input:/app/input ^
  -v $(pwd)/output:/app/output ^
  connecting-dots --mode 1B ^
  --persona "Investment Analyst" ^
  --job "Analyze revenue trends"
```

The container loads all PDFs, embeds sections, and writes `persona_ranking.json` to `output/` within 60 s.

## 7. Combined workflow

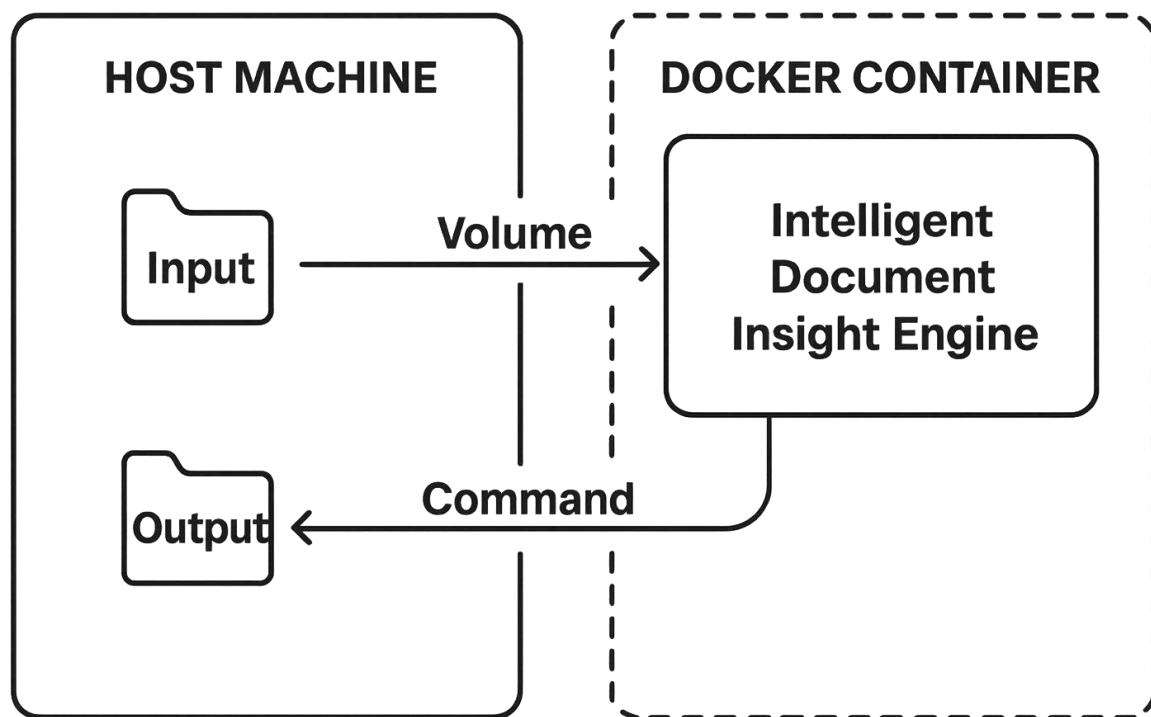
If you want both rounds in one shot:

```
docker run -- rm ^
  -v $(pwd)/input:/app/input ^
  -v $(pwd)/output:/app/output ^
  connecting-dots --mode both ^
  --persona "Product Manager" ^
  --job "Find UX best-practice sections"
```

The script first produces outline JSONs then immediately performs ranking using the extracted sections.

## 8. Understanding volume mounts

The `-v hostDir:containerDir` flags bind local folders so the engine can read PDFs and write results without rebuilding the image.



Volume mounts when running the Intelligent Document Insight Engine Docker container.

## 9. Troubleshooting

Symptom	Likely cause	Fix
"no such file or directory" when mounting	Windows path not shared with Docker	Check Docker Desktop → Settings → Resources → File Sharing
Processing takes > 60 s	Huge PDFs or > 10 files	Split PDFs or process batches
"SentenceTransformer import" error	Host lacks AVX and Docker falls back to older CPU mode	Set environment variable <code>FAST_EMBEDDINGS=1</code> to use TF IDF fallback

## 10. Optional: interactive FastAPI UI

The repo also contains a minimal FastAPI interface (`web_app.py`). To explore:

```
docker compose --profile web up --build
```

Then browse to `http://localhost:8000`. The UI lets you upload PDFs and trigger ranking from the browser.

## 11. Clean-up

Stop and remove all containers:

```
docker compose down
```

Remove the image if no longer needed:

```
docker rmi connecting-dots
```

## You're all set

With Docker handling dependencies and isolation, running the Intelligent Document Insight Engine is reduced to **three commands**: *build* → *mount* → *run*. Enjoy extracting smart outlines and persona-tailored insights from your PDFs!

\*  
\*\*